
2D Orthogonal Range Search

Mads Ravn, 20071580

Master's Thesis, Computer Science

January 2015

Advisor: Kasper Green Larsen



AARHUS
UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE

Abstract

► in English... ◄

Resumé

► in Danish . . . ◄

Acknowledgements



*Mads Ravn,
Aarhus, January 12, 2015.*

Contents

Abstract	iii
Resumé	v
Acknowledgments	vii
1 Introduction	3
I Theory	5
2 Related Work	7
2.1 kd-trees	7
2.2 Orthogonal Range Searching	7
2.2.1 Solving the ball-inheritance problem	8
2.2.2 Solving range reporting	9
2.3 Simplified Range Searching	11
3 ►...◄	15
4 Conclusion	17
Primary Bibliography	17

Chapter 1

Introduction

►Describe primary work◄ In an age where everybody is presented with the possibility of searching through a lot of data, for example ebay.com and bilzonen.dk, we are interested in fast queries with as little overhead space usage as possible. We want to do two-dimensional orthogonal range search in our database, which means we can select two axes on which we can select results between two points on each. We are essentially forming an axis-aligned rectangle and selecting all points which lie within it. ►Rewrite - Ellis style◄

Orthogonal Range Searching. *Orthogonal range searching* is one of the most fundamental and well-studied problems in computational geometry. Even with extensive research over three decades a lot of questions remain. In this thesis we will focus on 2D orthogonal range searching: Given n points from \mathbb{R}^2 we want to insert them into a data structure which will be able to efficiently report which k points lie within a given query range $Q \subseteq \mathbb{R}^2$. This query can be defined as two corners of a rectangle, the lower left corner and the upper right corner, seeing as the query range is orthogonal to the axes.

RAM, I/O, pointer models ►Skriv om Word RAM Model sammenlignet med de andre typer her. Når vi pakker bits ned i et word kan vi stadig tilgå det hele da access er constant time og operations er constant time◄

Rank Space Reduction. Given n points from a universe U , the rank of a given point is defined as the amount of points which precede it in a sorted list. Given two points $a, b \in U : a < b$ iff $rank(a) < rank(b)$. Expanding this concept to 2 dimensions we have a set P of n points on a $U \times U$ grid. We compute for the x -rank r_x for each point in P by finding the rank of the x-coordinate amongst all the x-coordinates in P . The y -rank r_y finds the rank of y-coordinate amongst all of the y-coordinates in P . Using *rank space reduction* on P a new set P^* is constructed where $(x, y) \in P$ is replaced by $(r_x(x), r_y(y)) \in P^*$. Given a range query $q = [x_1, x_2] \times [y_1, y_2]$, a point $(x, y) \in P$ is found within q iff $(r_x(x), r_y(y))$ is found within $q^* = [r_x(x_1), r_x(x_2)] \times [r_y(y_1), r_y(y_2)]$. ►Noget med at deres ordered property remains intact.◄ Computing the set P^*

from P using rank space reduction, P^* is said to be in rank space. While the n points could be represented by $\lg U$ bits in P , they can now be represented by $\lg n$ bits in P^* with $\lg n \ll \lg U$ which saves memory. ►We have essentially created a mapping between ...◄

Ball Inheritance. Given a perfect binary tree with n leaves and n labelled balls at the root the idea is to distribute the balls from the root to the leaves. The balls at the root are contained in an ordered list and for each ball in a nodes list one of its children is picked to inherit the ball such that both children receive the same amount of balls. Each level of the tree contains the same amount of balls, and at level i each node contains 2^i balls. Eventually each ball reaches a leaf of the tree and each leaf will contain exactly one ball. The identity of a given ball is a node and the index of the ball in that nodes list. The goal is to track a balls inheritance from a given node to a leaf and report the identity of the leaf. When referring to ball inheritance, the words *ball* or *point* can be used interchangeably.

Outline.

Notation. The set of integers $\{i, i + 1, \dots, j - 1, j\}$ is denoted by $[i, j]$. When no base is explicitly given logarithm will have base 2. ϵ is an arbitrary small constant. Given an array A , $A[i]$ denotes the entry with index i in A and $A[i, j]$ denotes the subarray containing the entries from i to j in A . $A[1..n]$ denotes an array A of size n with entries 1 to n . Throughout the thesis the successor of x will be meant as the first number which is greater or equal to x - the same applies for predecessor of x . The work will be done under the assumption that no two points will have the same x-coordinate and no two points will have the same y-coordinate. ►Må vi lave den antagelse hvis vi beskriver hvorfor og hvordan den kan “rettes” senere?◄

Part I

Theory

Chapter 2

Related Work

In this chapter the original work of Chan et al. [2] is presented followed by a simplification of their idea. This simplification serves two purposes: It will be easier to implement which will present a much cleaner code to read and it will be easier to execute since it has much less things to compute and look up. The theory behind a kD -tree will also be explained as it is the de facto standard of orthogonal range queries today and will be used to compare the practical results of the simplified model.

The algorithms mentioned in this chapter are all *output-sensitive*, meaning that their running time depends on the amount of results found.

2.1 kd-trees

The current standard of range reporting is kd-trees. This technique will be used as a reference point when evaluating the results of the primary work in this thesis. With linear space it is optimal for range reporting on the RAM.

A kd-tree is constructed recursively: Given n points, the median of the points with respect to x are found. All points which has an x -coordinate larger than the median goes to the right child, while points which has an x -coordinate smaller than the median, and the median point, goes to the left child. At the next level the points will be divided in a similar fashion, this time using the y -median and the y -coordinates instead. When dividing n points, the median will be chosen as the $\lceil n/2 \rceil$ -th smallest number [1]. Switching back and forth between focussing on x or y is done at each level of the tree and this will be continued until a binary tree has been constructed with n leaves.

The tree with n points can be represented as flat array with n entries. The $\lceil n/2 \rceil$ -th element in the array is the root of the tree. With this in mind, the entire tree can be represented using a flat array with n elements. **►Far from done - section needs work and rework◄**

2.2 Orthogonal Range Searching

Utilizing the *ball-inheritance* concept, Chan et al. propose a better solution for orthogonal range search queries. With the same space complexity as the kD -

tree, but lower query time. Theorem 2.1 states that *for any $2 \leq B \leq \lg^\epsilon n$, we can solve 2-d orthogonal range reporting in rank space with $\mathcal{O}(n \lg_B \lg n)$ space and $(1+k)\mathcal{O}(B \lg \lg n)$ query time.*

First we will look at how ball distributing will work in the ball-inheritance tree and how the ball-inheritance problem can be solved. Then we will show how the ball-inheritance is used in conjunction with other data structures to find the points within a search query $q = [x_1, x_2] \times [y_1, y_2]$.

2.2.1 Solving the ball-inheritance problem

Consider a perfect binary tree with n leaves. At each level a bit vector $A[1..n]$ is used to indicate which of a node's children a ball is inherited by: If $A[i]$ is 0 means that the ball with identity i in that node was inherited by its left child and 1 means that it was inherited by its right child. Given a node and an identity of a ball we can know calculate the ball's identity in the node it is inherited by. The node can answer the query $\text{rank}(k) = \sum_{i \leq k} A[i]$. If a ball is inherited by the right child node its new identity at that node is $\text{rank}(i)$ because that is how many 1's that precede it in the current node. If a ball is inherited by the left child node the new identity is then $i - \text{rank}(i)$. With this information it is possible to traverse down the tree following a ball from any given node to a leaf. There are n balls per level which is represented by a bit vector of size n bits per level. Even though conceptually this bit vector is divided out amongst the nodes of that level, we can interchangeably think of a bit vector per level or a bit vector per node. **►Rephrase and replace◄** Each level in the tree uses $\mathcal{O}(n)$ bits to store the bit vectors. This adds up to $\mathcal{O}(n \lg n)$ bits, or $\mathcal{O}(n)$ words in all. This trivial solution to the ball-inheritance problem uses $\mathcal{O}(\lg n)$ query time, given that it follows a ball $\mathcal{O}(\lg n)$ steps down to its leaf. The rank function is a constant time query.

A bit vector is an array with entries from the alphabet $\Sigma = \{0, 1\}$, where each entry is used to indicate whether a left or right child has been chosen to inherit a given ball. By expanding the alphabet we can point to the childrens children, $\Sigma = \{0, 1, 2, 3\}$, the childrens childrens children, $\Sigma = \{0, 1, 2, 3, 4, 5, 6, 7\}$, and so forth. Expanding the alphabet will use $\mathcal{O}(n \lg \Sigma)$ bits per level. Storing a pointer from level i to level $i + \Delta$ increases the storage space by Δ bits per ball, but also enables the ball to be inherited by 2^Δ descendants. By expanding the alphabet the query time can be lowered since it is possible to take bigger steps down the tree.

Using this concept, we pick B such that $2 \leq B \leq m$, where m is the height of the tree. All levels that are a multiple of B^i expand their alphabet such that the balls reach B^i levels down. If a target level does not exist, the ball points to its leaf. We need at most visit B levels that are multiple of B^i before reaching a level that is multiple of B^{i+1} , making it possible to jump down the tree with bigger and bigger steps.

Storing the expanded alphabets at each level that is multiple of B^i costs B^i bits per ball. The total cost is then $\sum_i \frac{\lg n}{B^i} \cdot \mathcal{O}(B^i) = \mathcal{O}(\lg n \cdot \lg_B \lg n)$ bits per ball, at all levels. With n balls, this is $\mathcal{O}(n \lg_B \lg n)$ words of space with query

time of $\mathcal{O}(B \lg_B \lg n)$. ►Spørg Kasper med matematikken bag dette - især om “at all levels” er rigtig og hvordan Σ forvandles til $\lg_B \lg n$ ◄

2.2.2 Solving range reporting

With a solution to the ball-inheritance problem, Chan et al. [2] proposes **Lemma 2.4** *if the ball inheritance problem can be solved with space S and query time τ , 2-d range reporting can be solved with space $\mathcal{O}(S + n)$ and query time $\mathcal{O}(\lg \lg n + (1 + k)\tau)$.*

Consider a perfect binary tree with n leaves. The root contains n points in 2-d rank space. These n balls are sorted by their y-rank and will keep this order in all of the nodes in the tree. When distributing the balls for inheritance, a node will give both its children half of its balls: the lower half sorted by the x-rank to its left child and the upper half by x-rank to its right child. The actual coordinates of the balls are only stored at the leaves and we know that x-coordinates of the balls in the leaves are sorted from left to right - smallest to highest. Since the actual coordinate points are only stored once, this data structure uses linear space.

Having distributed the n points from the root to the leaves, additional data structures are required in order to answer the range queries. A succinct data structure will be used for the *range minimum query*. Consider an array A with n comparable keys, this data structure allows finding the index of the minimum key in the subarray $A[i, j]$. ►Describe either the cartesian tree or reference 29◄. For each node in the tree that is a right child a range minimum query structure is added. The indices are the y-rank and the keys are the x-rank that the given node contains. A range maximum query structure is added to the all the nodes which are left children. Each data structure uses $2n + \mathcal{O}(n)$ bits, making it $\mathcal{O}(n)$ bits per level of the tree and $\mathcal{O}(n \lg n)$ bits in all - i.e. $\mathcal{O}(n)$ words of space. This range minimum data structure has a constant query time. ►Skal jeg inkludere samme reference som i Chan et al. [2] [29]?◄

There is still the matter of adding predecessor (and successor) search for the y-rank. For this purpose a succinct data structure is added to some levels of the tree. Given a sorted array $A[1..n]$ of ω -bit integers, predecessor search queries in $\mathcal{O}(\lg \omega)$ time is supported using $\mathcal{O}(n \lg \omega)$ space which has oracle access to the entries in the array. ►reference◄. The points in rank space of $\mathcal{O}(\lg n)$ bits will use $\mathcal{O}(n \lg \lg n)$ bits per level, with $\omega = \lg n$, and $\mathcal{O}(n \lg n \lg \lg n)$ bits in all, which is $\mathcal{O}(n \lg \lg n)$ words. In order to reduce this to linear space we will only place this predecessor search structure at levels which are multiples of $\lg \lg n$. When using the predecessor search from the lowest common ancestor of x_1^* and x_2^* , $LCA(x_1^*, x_2^*)$, we go up to the closest ancestor to use its search structure. Searching takes $\mathcal{O}(\lg \lg n)$ time plus $\mathcal{O}(1)$ queries to the ball-inheritance structure: using the ball-inheritance structure we walk at most $\lg \lg n$ steps down while translating the ranks of y_1 and y_2 to the right and left child of $LCA(x_1^*, x_2^*)$.

The reason why this structure is necessary for the y-ranks and not the x-ranks, is because of the way the points have been distributed in the ball-

inheritance tree: From left to right, the leaves have x-rank $1, 2, \dots, n$ so we can easily locate a given range in x-dimension, but in order to keep track of the y-dimensional range we need to follow the balls down the ball-inheritance structure. Adding this structure to each $\lg \lg n$ level saves us from going all the way from the root down to the *LCA*. **►rephrase◄**

In order to use this new data structure to report points in the range of $[x_1, x_2] \times [y_1, y_2]$ we follow these steps:

1. We find the rank space successor of x_1 , x_1^* , and the rank space predecessor of x_2 , x_2^* . We use these to find the lowest common ancestor of x_1^* and x_2^* , $LCA(x_1^*, x_2^*)$. This is the lowest node in the tree containing at least all the points between x_1 and x_2 . By knowing x_1^* and x_2^* , finding the lowest common ancestor is a constant time operation. Given that points are in an array, we can use xor as our *LCA* operation. **►rephrase◄**
2. As in step 1 where we found the rank space of the x-coordinates, we find the rank space coordinates of the y-coordinates, y_1^* and y_2^* , inside the left and right child of $LCA(x_1^*, x_2^*)$. This step is precisely what the succinct data structure mentioned above supports.
3. We now descend into the right child of $LCA(x_1^*, x_2^*)$ and use the range minimum query structure to find the index m (the y-rank) of the point with the smallest x-rank in the range $[y_1^*, y_2^*]$. Solving the ball-inheritance problem, we follow the path to the leaf to find the actual x-coordinate of the point. If the x-coordinate is smaller than x_2 we return the point as a result and recurse into the ranges of $[y_1^*, m - 1]$ and $[m + 1, y_2^*]$ in order to find more points. When this is done we apply the same concept to the left child of $LCA(x_1^*, x_2^*)$ using the range maximum query to find points above x_1 .

►Insert figure to conceptually show we are working our way out from the inside◄

The time complexity of step 3 depends on the use of the ball-inheritance structure. The time to traverse this structure is dependent on the improvements made in 2.2.1. An empty range will result in two queries, one query to each child of $LCA(x_1^*, x_2^*)$, otherwise it will not exceed twice the number of nodes beneath $LCA(x_1^*, x_2^*)$. **►Why?◄** Conceptually, $LCA(x_1^*, x_2^*)$ describes a point between x_1 and x_2 . Step 3 selects points that are in the range of $[y_1, y_2]$ moving outwards from the point of $LCA(x_1^*, x_2^*)$, always picking the point closest to $LCA(x_1^*, x_2^*)$ in its decreasing y-range. **►rephrase◄**

Going back to Lemma 2.4, we see that the time complexity fits: $\mathcal{O}(\lg \lg n)$ time is used for the predecessor search and $\mathcal{O}((1 + k)\tau)$ time is used for walking from the *LCA* to the leaves solving the ball inheritance problem for the k results.

2.3 Simplified Range Searching

In this section we will describe a simplified model of the one mentioned in section 2.2. It mainly uses the ball-inheritance structure. A predecessor search structure will be used at the root of the tree to find the rank space of two points, (x_1, y_1) and (x_2, y_2) , delimiting the search range. The algorithm will be described under the assumption that the solution to the ball inheritance problem requires $\mathcal{O}(\lg n)$ query time. **►Nu hvor jeg lige har beskrevet den originale, hvor meget skal jeg så gå i detaljer med hvorfor vi gør tingene? Som f.eks. at have tingene i rank space◄**

The ball distribution works the same way as it did in the original range searching by Chan et al.: Given a perfect tree with n leaves, the n points are sorted by their y-rank at the root. Distributing its balls, a node will give half of its points to each of its children: the smallest by x-rank to its left child and the highest by x-rank to its right child. The points in the two new lists will retain the same order as the parent.

Given a range query $q = [x_1, x_2] \times [y_1, y_2]$ we are going to find the rank successors of x_1 and y_1 and the rank predecessors of x_2 and y_2 . **►Remember to explain why◄** We call these x_1^*, y_1^*, x_2^* and y_2^* . We use x_1^* and x_2^* to find the lowest common ancestor, $LCA(x_1^*, x_2^*)$, containing all the points between x_1 and x_2 . We know the positions of the leaves containing x_1^* and x_2^* so we can traverse from the LCA down to them both. Traversing to x_1^* , the first stop is the left child of the LCA . From here, each time a node selects its left child as the path to x_1^* we know that the subtree contained in the right child contains points between x_1 and x_2 . Symmetrically, the same applies when going right from the LCA : Each time a node selects a right child on the path to x_2^* the subtree contained in the left child contains points between x_1 and x_2 . This concept is seen on figure 2.2.

In the root we mark the position of y_1^* and y_2^* on the bit vector used to indicate which descendant a balls goes to. This range indicates which balls lie within the range of $[y_1, y_2]$. Going forward we will use i and j to indicate this range in the bit vector of any given node. When a node inherits this range from a parent, the rank function is used to query how many points before i went to this node, and how many points before j went to this node. We refer to section 2.2.1 on how to update the positions using the rank query. This works because the points in a node are sorted by their y-rank and we keep track of how many points in a given node falls within the range of $[y_1, y_2]$ updating the range of $[i, j]$ everytime a node inherits balls.

While keeping the y-rank range updated, we travel from the root to the LCA . From here we travel down to all the leafs inside the range of $[x_1, x_2]$ while still updating the y-range. Using the y-range we might be able to decide if we can fully include or fully exclude entire subtrees. We can fully include a subtree if the entire y-rank range of the subtree is within the $[y_1, y_2]$ or fully exclude it if the entire y-range of the subtree falls without $[y_1, y_2]$. Referring to

figure 2.1 we see how we conceptually update the bit vector and how it would be obvious if the range is either empty or contain the entire subtree. Generalizing this, a leaf can be treated as a subtree: either including or excluding it.

Summing up, the data structure is utilized as follows:

1. Use a binary search to find the rank space predecessors and successors of x_1, y_1, x_2 and y_2 . At this point the algorithm will terminate if either rank space range is empty.
2. From $LCA(x_1^*, x_2^*)$, each leaf between x_1^* and x_2^* will be visited. The range y -rank range found in step 1 will be updated from the root to this LCA and updated from the LCA to all the leaves visited. When a leaf is visited, its actual coordinates will be reported back as a result.

In order to make the rank query a constant time query, we create a partial rank sum list of the bit vector. Given a bit vector with n bits, the rank at certain intervals is precalculated and stored. Storing the rank at each 32 bits then enable us to calculate $rank(i)$ by looking up entry $\lfloor \frac{i}{32} \rfloor$ in the partial sum list and summing up $i \% 32$ bits from the bit vector. Calculating the rank thusly takes $\mathcal{O}(1)$ time and requires $\mathcal{O}(n)$ words of space. **►Rephrase and maybe move to other chapter◄**

The actual coordinates of the points are only stored at the leaves which then takes up $\mathcal{O}(n)$ words of space. The rest of the tree contains $\lg n$ levels of bit vectors of n bits taking $\mathcal{O}(n \lg n)$ bits, $\mathcal{O}(n)$ words. Looking up the rank-space predecessor and successor of x_1, x_2, y_1 and y_2 using a simple binary search requires $\mathcal{O}(n)$ space and $\mathcal{O}(\lg n)$ time. Summing it up, the entire data structure uses $\mathcal{O}(n)$ words of space.

Walking from the root to the LCA requires $\lg n$ steps. Visiting each of the k leaves containing the points which will be reported as a result takes $\mathcal{O}(k \cdot \lg n)$ time.

This adds up to $\mathcal{O}(\lg n + k \cdot \lg n)$ query time to report k points as results. An empty range will be detected by the binary search. If the binary search does not report an empty range, we proceed to use the algorithm, visiting each leaf between x_1^* and x_2^* taking $\mathcal{O}(\lg n)$ time to visit each leaf.

The three approaches described above all use a number of bits linear to the amount of points. The original search structure by Chan et al. [2] is the fastest with its $(1+k)\mathcal{O}(B \lg \lg n)$ query while kd-trees are the slowest with $\mathcal{O}(\sqrt{n} + k)$ query time.




Figure 2.1: When nodes inherit their bit vector ranges from their parent, it can become obvious if the entire subtree is contained within the range of $[y_1, y_2]$ or falls without the range.




Figure 2.2: Traversing left from the LCA, each right subtree contains x-coordinates between x_1 and x_2 . Traversing right from the LCA the same holds for left subtrees.

Chapter 3



Chapter 4

Conclusion



Bibliography

- [1] Mark de Berg, Otfried Cheong, Marc van Kreveld, and Mark Overmars. *Computational Geometry: Algorithms and Applications*. Springer-Verlag TELOS, Santa Clara, CA, USA, 3rd ed. edition, 2008. ISBN 3540779736, 9783540779735.
- [2] Timothy M. Chan, Kasper Green Larsen, and Mihai Pătraşcu. Orthogonal range searching on the ram, revisited. In *Proceedings of the Twenty-seventh Annual Symposium on Computational Geometry*, SoCG '11, pages 1–10, New York, NY, USA, 2011. ACM. ISBN 978-1-4503-0682-9. doi: 10.1145/1998196.1998198. URL <http://doi.acm.org/10.1145/1998196.1998198>.